Word Embedding based Search Engine

Brijesh Dabhi
brijesh.dabhi@iitgn.ac.in
Indian Institute of Technology Gandhinagar

Dhruv Goel
23110098@iitgn.ac.in
Indian Institute of Technology Gandhinagar

Rahul Singal
rahul.singal@iitgn.ac.in
Indian Institute of Technology Gandhinagar

Rupak Banerjee
rupak.banerjee@iitgn.ac.in
Indian Institute of Technology Gandhinagar

Harshit
harshit.harshit@iitgn.ac.in
Indian Institute of Technology Gandhinagar

Abstract

Traditional search systems rely on keyword matching techniques such as TF-IDF and BM25, which fail to capture semantic relationships between queries and documents. This project implements a semantic search engine built on the AG News dataset, using Word2Vec and Sentence-BERT as embedding backends with Approximate Nearest Neighbor (ANN) retrieval via PyNNDescent. We further augment query representations with nearest-term query expansion. The system is evaluated on 1,000 sampled test queries against an indexed corpus of 50,000 documents using Precision@k, Recall@k, MAP, and nDCG@k. With query expansion, Word2Vec achieves Precision@5 of 0.7510 and MAP of 0.6464, outperforming the no-expansion baseline across every metric. Results confirm that embedding-based retrieval captures semantic similarity effectively, and that query expansion provides a consistent and measurable ranking improvement.

1 Introduction

The exponential growth of textual data has made efficient and accurate information retrieval a critical challenge. Traditional search systems based on keyword matching — such as TF-IDF and BM25 — fail to capture semantic relationships between words, leading to poor retrieval when queries and relevant documents share meaning but differ in vocabulary.

Recent advances in Natural Language Processing (NLP), particularly word embeddings such as Word2Vec and BERT, represent text in a continuous vector space where semantic similarity is preserved. This project builds a complete retrieval pipeline on the AG News benchmark corpus, going beyond the basic assignment requirements by implementing two backends, adding nearest-term query expansion, using Approximate Nearest Neighbor (ANN) search for scalable retrieval, and extending evaluation to include MAP and nDCG in addition to the required Precision@k and Recall@k.

2 Problem Statement

Despite advances in search systems, the following core challenges remain:

- Inability to capture semantic similarity between queries and documents when vocabulary differs.
- Poor retrieval for paraphrased or synonym-heavy queries.
- High computational cost for brute-force nearest neighbor search at scale.
- Limited evaluation coverage when only precision and recall are reported.

This project addresses these limitations using dense embedding representations, ANN-based retrieval, and query expansion, all evaluated with a comprehensive suite of IR metrics.

3 Dataset

We use the AG News dataset [1], a widely-used benchmark for topic classification and information retrieval.

- **Train split:** 50,000 documents indexed as the retrieval corpus.
- **Test split:** 7,600 documents, from which 1,000 are sampled as evaluation queries.
- **Classes:** Four topic categories — World (1), Sports (2), Business (3), and Sci/Tech (4).

Relevance definition. A retrieved document is treated as relevant if its class label matches the query's class label. This is the standard IR proxy protocol for AG News and is simple, automated, and reproducible.

An important consequence of this definition is that Recall@k values are inherently very small: each class contains approximately 12,500 relevant documents in the train split, so top-k retrieval for small k can only surface a tiny fraction. For this reason, MAP and nDCG are our primary quality signals.

4 System Architecture

The end-to-end retrieval pipeline consists of six sequential stages:

(1) **Load:** AG News train and test splits are loaded and split by class.
(2) **Embed:** Every document is encoded into a dense vector using the selected backend (Word2Vec or BERT).
(3) **Index:** An ANN index is built over the document vectors using PyNNDescent, with an exact cosine fallback for small or fragile runs.
(4) **Expand:** For Word2Vec mode, each query token's nearest neighbors in term space are appended to the query before encoding.
(5) **Search:** Top-k documents are retrieved by cosine similarity.
(6) **Evaluate:** Precision@k, Recall@k, MAP, and nDCG@k are computed over all evaluated queries.

All document and query vectors are L2-normalised so that cosine similarity reduces to a dot product, which is faster and numerically stable.

5 Embedding Backends

5.1 Word2Vec

We train a skip-gram Word2Vec model from scratch on the AG News training corpus using Gensim:

- **Algorithm:** Skip-gram
- **Dimensions:** 200
- **Window size:** 5
- **Minimum token count:** 2
- **Training epochs:** 2,000

Document vectors are computed as the mean of their token vectors. Tokens below the minimum count threshold are excluded at training time and skipped at inference time.

5.2 Sentence-BERT

We use the pretrained `sentence-transformers/all-MiniLM-L6-v2` model:

- **Dimensions:** 384
- **Training:** No fine-tuning; inference only.
- **Encoding:** Full sentence embeddings via mean pooling of token representations.

Both backends expose the same fit/search/expand API, so the backend can be swapped with a single CLI flag.

6 Approximate Nearest Neighbor Retrieval

Brute-force retrieval requires $O(N \cdot d)$ comparisons per query, which is prohibitive at scale. We use PyNNDescent [6], which builds a k-nearest-neighbor graph offline and traverses it at query time, achieving sub-linear query cost in practice.

- **Index construction:** $O(N \log N)$ offline.
- **Query time:** Approximately $O(\log N)$.
- **Fallback:** If the ANN index fails to initialise (e.g., on very small subsets), the engine automatically falls back to exact cosine search with a warning, so no run crashes.

7 Query Expansion

Short news headlines are often noisy and may miss vocabulary present in relevant documents. To mitigate this, we implement nearest-term query expansion for Word2Vec mode:

(1) For each token in the query, retrieve its k nearest neighbors in the Word2Vec term vector space.
(2) Append these neighbor terms to the query token list.
(3) Re-encode the augmented query as the mean of all token vectors (original + expanded).

This shifts the query vector toward the semantic neighborhood of the query tokens, effectively acting as a soft thesaurus. Table 1 shows a concrete example.

A known limitation is that Word2Vec co-occurrence neighbors are occasionally noisy. For example, the query token "Intel" produced expanded terms such as "plantation" and "libraries." Expansion is net-positive overall, but can occasionally drag individual query results off-topic.

Table 1: Query expansion example. The original query tokens are augmented with semantically related terms before encoding.

Field	Content
Original query	*Israeli Soldiers Kill Palestinian Gunman in Gaza*
Expanded terms	armored, jenin, nablus, mosul, wound, infiltrate, Abbas
Top result	*Israeli Army Kills Gunman in Gaza Strip* (score: 0.91)

8 Evaluation Metrics

We report four standard IR metrics over 1,000 evaluation queries at $k \in \{5, 10, 20\}$:

- **Precision@k:** Fraction of the top-k retrieved documents that are relevant.
- **Recall@k:** Fraction of all relevant documents in the corpus that appear in the top-k.
- **MAP:** Mean Average Precision; rewards placing relevant documents early in the ranked list.
- **nDCG@k:** Normalized Discounted Cumulative Gain; discounts relevant hits at lower ranks, closely approximating perceived user quality.

All averages are computed over the full 1,000 evaluated queries (zero queries were skipped in any run).

9 Results

9.1 Quantitative Results

Table 2 presents the complete metric breakdown for Word2Vec without and with query expansion across all k values.

9.2 Qualitative Results

To verify semantic quality beyond aggregate metrics, we inspect the engine's output on individual queries.

Example: "Intel calls for Internet overhaul" (Sci/Tech, class 4). The top results included: *Intel suggests upgrading internet to Intelnet* (0.79), *Intel Proposes Internet Transformation* (0.79), *Intel plans to fix internet* (0.78), and *IDF – Intel seeks role in development of the next Internet* (0.77). Seven of the top eight results are correctly classified as Sci/Tech. The single off-class hit was a business article about internet infrastructure investment — topically adjacent and not a spurious match.

Example: "CRM Best Practices: TCO and ROI" (Sci/Tech, class 4). The top-5 results are all CRM-related Sci/Tech articles, with cosine scores between 0.74 and 0.82. The engine correctly identifies that "CRM," "TCO," and "ROI" are business software concepts, retrieving articles about Siebel, Oracle CRM, and SugarCRM rather than generic business or finance content.

10 Analysis

10.1 Effect of Query Expansion

Query expansion improves every precision and ranking metric across all k values. The MAP gain of +0.033 is the largest absolute improvement, indicating that expansion pushes relevant documents higher in the ranking — not merely that more of them appear in

Table 2: Word2Vec retrieval results on AG News (50k train, 1k queries). Best row highlighted.

Mode	P@5	P@10	P@20	R@5	R@10	R@20	MAP	nDCG@5	nDCG@10	nDCG@20
Word2Vec (baseline)	0.7386	0.7197	0.7055	0.0003	0.0006	0.0011	0.6132	0.7419	0.7273	0.7146
Word2Vec + expansion	**0.7510**	**0.7389**	**0.7252**	**0.0003**	**0.0006**	**0.0012**	**0.6464**	**0.7529**	**0.7438**	**0.7324**
Δ from expansion	+0.0124	+0.0192	+0.0197	≈0.000	≈0.000	+0.001	+0.0332	+0.0110	+0.0165	+0.0178

the top-k. The improvement is consistent across all 1,000 queries, ruling out overfitting to a small test subset.

10.2 Interpretation of Low Recall

The low Recall@k values (≈ 0.0003–0.0012) are an artifact of the evaluation protocol rather than a failure of the retrieval engine. AG News has approximately 12,500 relevant documents per class in the 50k training corpus. Top-20 retrieval can cover at most 20 of them, so recall is bounded above at $20/12500 \approx 0.0016$ regardless of engine quality. MAP and nDCG are therefore the appropriate primary quality metrics for this evaluation setup.

10.3 Word2Vec vs. BERT Trade-off

Word2Vec (averaged skip-gram, 200d) achieves P@5 of 0.75 on a topical corpus like news, where topic categories map cleanly to vocabulary clusters. Sentence-BERT typically achieves stronger semantic generalization but at significantly higher inference cost (no training required, but each document requires a full transformer forward pass). For latency-constrained deployments on topical corpora, Word2Vec offers a favorable accuracy/cost trade-off.

10.4 Limitations

- **Coarse relevance labels:** Class-label matching treats all within-class documents as equally relevant, which is a rough approximation of true semantic relevance.
- **Noisy expansion:** Some Word2Vec nearest neighbors are co-occurrence artefacts rather than semantic synonyms. Weighting expansion terms by their embedding distance could mitigate this.
- **BERT evaluation at scale:** Full BERT evaluation at 50k documents and 1,000 queries was not run due to CPU inference cost; a direct comparison at equivalent scale remains future work.

11 Engineering and Reproducibility

The implementation is organized as a modular Python package under `src/search_engine/`:

- `data.py` — AG News loader and train/test split inference.
- `embeddings.py` — Word2Vec and BERT embedding logic with L2 normalization.
- `ann.py` — PyNNDescent ANN wrapper with exact-search fallback.
- `engine.py` — Search engine fit, search, and query expansion logic.
- `evaluation.py` — Metric computation and query accounting.
- `text_utils.py` — Tokenizer utilities.

Three CLI scripts are provided: `run_experiment.py` for end-to-end runs, `train_and_save_models.py` for per-epoch checkpoint saving, and `evaluate_saved_model.py` for reloading and re-evaluating checkpoints. A minimal Streamlit app (`app.py`) supports live interactive querying with configurable backend and expansion toggle.

- **Modules:** 9 Python modules, ≈1,500 lines of code.
- **Checkpoints:** 148 MB total (Word2Vec epoch 2000: 72 MB; BERT epoch 2000: 76 MB).
- **Reproducibility:** Controlled via `-seed`, `-w2v-seed`, and `-ann-random-state` flags.
- **One-command reproduction:** A full experiment can be reproduced with a single CLI invocation.

Listing 1: One-command experiment reproduction.

```
python scripts/run_experiment.py \
    --dataset sh0416/ag_news \
    --embedding word2vec \
    --max-train 50000 \
    --eval-queries 1000 \
    --k 5 10 20
```

12 Applications

The semantic retrieval approach demonstrated in this project is applicable to a range of real-world systems:

- News and document search engines.
- Question answering and reading comprehension systems.
- Recommendation engines based on content similarity.
- Chatbot response retrieval.
- Enterprise knowledge base search.

13 Conclusion

This project delivers a complete semantic search engine over the AG News corpus with two embedding backends (Word2Vec and Sentence-BERT), ANN-based retrieval, nearest-term query expansion, and an extended IR evaluation suite. The Word2Vec backend with query expansion achieves Precision@5 of 0.75 and MAP of 0.65 on 1,000 test queries against a 50k-document corpus. Query expansion consistently improves precision and ranking metrics by 1–4% across all k values. The codebase is modular, reproducible, and ships a working interactive query application. The study demonstrates that simple averaged skip-gram embeddings are highly competitive on topical corpora, and that nearest-neighbor query expansion is a practical and effective technique for improving retrieval quality without retraining.

Acknowledgments

We thank Dr. Anirban and the CS 328 teaching team for the project brief, support and guidance.

References

[1] Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level Convolutional Networks for Text Classification. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2015.

[2] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. Distributed Representations of Words and Phrases and their Compositionality. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2013.

[3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of NAACL-HLT*, 2019.

[4] Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of EMNLP-IJCNLP*, 2019.

[5] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-Scale Similarity Search with GPUs. *IEEE Transactions on Big Data*, 2019.

[6] Leland McInnes, John Healy, and James Melville. UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction. *arXiv preprint arXiv:1802.03426*, 2018.

[7] Stephen Robertson and Hugo Zaragoza. The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval*, 3(4):333–389, 2009.